


# CS 188: Artificial Intelligence

## Informed Search



Fall 2022

University of California, Berkeley

# Today


---

- Informed Search
  - Heuristics
  - Greedy Search
  - A\* Search
- Graph Search



# Recap: Search

---



# Recap: Search

- Search problem:


- States (configurations of the world)
- Actions and costs
- Successor function (world dynamics)
- Start state and goal test

- Search tree:

- Nodes: represent plans for reaching states
- Plans have costs (sum of action costs)


- Search algorithm:

- Systematically builds a search tree
- Chooses an ordering of the fringe (unexplored nodes)
- Optimal: finds least-cost plans



# Example: Pancake Problem

---



Cost: Number of pancakes flipped

# Example: Pancake Problem

## **BOUNDS FOR SORTING BY PREFIX REVERSAL**

William H. GATES

*Microsoft, Albuquerque, New Mexico*

Christos H. PAPADIMITRIOU\*†

*Department of Electrical Engineering, University of California, Berkeley, CA 94720, U.S.A.*

Received 18 January 1978


Revised 28 August 1978

For a permutation  $\sigma$  of the integers from 1 to  $n$ , let  $f(\sigma)$  be the smallest number of prefix reversals that will transform  $\sigma$  to the identity permutation, and let  $f(n)$  be the largest such  $f(\sigma)$  for all  $\sigma$  in (the symmetric group)  $S_n$ . We show that  $f(n) \leq (5n+5)/3$ , and that  $f(n) \geq 17n/16$  for  $n$  a multiple of 16. If, furthermore, each integer is required to participate in an even number of reversed prefixes, the corresponding function  $g(n)$  is shown to obey  $3n/2 - 1 \leq g(n) \leq 2n + 3$ .



# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```








# Informed Search

---




# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem
  - Examples: Manhattan distance, Euclidean distance for pathing




# Example: Heuristic Function



$h(x)$

# Example: Heuristic Function

Heuristic: the number of the largest pancake that is still out of place




# Greedy Search





# Greedy Search

- Expand the node that seems closest...





- What can go wrong?



# Greedy Search

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS



[Demo: contours greedy empty (L3D1)]

[Demo: contours greedy pacman small maze (L3D4)]

# Video of Demo Contours Greedy (Empty)

---





# Video of Demo Contours Greedy (Pacman Small Maze)

---




# A\* Search

---



# A\* Search

---



UCS





Greedy





A\*

# Uniform-Cost Search




# Greedy Search



# Combining UCS and Greedy



- **Uniform-cost** orders by path cost, or *backward cost*  $g(n)$
- **Greedy** orders by goal proximity, or *forward cost*  $h(n)$



- **A\* Search** orders by the sum:  $f(n) = g(n) + h(n)$

# Combining UCS and Greedy



- **Uniform-cost** orders by path cost, or *backward cost*  $g(n)$
- **Greedy** orders by goal proximity, or *forward cost*  $h(n)$



- **A\* Search** orders by the sum:  $f(n) = g(n) + h(n)$

# When should A\* terminate?



- Should we stop when we enqueue a goal?



- No: only stop when we dequeue a goal




# Is A\* Optimal?




- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!


# Admissible Heuristics



# Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs


# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

- Examples:




4



- Coming up with admissible heuristics is most of what's involved in using A\* in practice.

# Optimality of A\* Tree Search

---




# Optimality of A\* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- $h$  is admissible

Claim:


- A will exit the fringe before B



# Optimality of A\* Tree Search: Blocking

Proof:

- Imagine  $B$  is on the fringe
- Some ancestor  $n$  of  $A$  is on the fringe, too (maybe  $A$ !)
- Claim:  $n$  will be expanded before  $B$ 
  1.  $f(n)$  is less or equal to  $f(A)$



# Optimality of A\* Tree Search: Blocking

## 1. $f(n)$ is less than or equal to $f(A)$

- Definition of f-cost says:

$$f(n) = g(n) + h(n) = (\text{path cost to } n) + (\text{est. cost of } n \text{ to } A)$$

$$f(A) = g(A) + h(A) = (\text{path cost to } A) + (\text{est. cost of } A \text{ to } A)$$

- The admissible heuristic must underestimate the true cost

$$h(A) = (\text{est. cost of } A \text{ to } A) = 0$$

- So now, we have to compare:

$$f(n) = g(n) + h(n) = (\text{path cost to } n) + (\text{est. cost of } n \text{ to } A)$$


$$f(A) = g(A) = (\text{path cost to } A)$$

- $h(n)$  must be an underestimate of the true cost from  $n$  to  $A$

$$(\text{path cost to } n) + (\text{est. cost of } n \text{ to } A) \leq (\text{path cost to } A)$$

$$g(n) + h(n) \leq g(A)$$

$$f(n) \leq f(A)$$






# Optimality of A\* Tree Search: Blocking

Proof:

- Imagine  $B$  is on the fringe
- Some ancestor  $n$  of  $A$  is on the fringe, too (maybe  $A$ !)
- Claim:  $n$  will be expanded before  $B$ 
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$



# Optimality of A\* Tree Search: Blocking

## 2. $f(A)$ is less than $f(B)$

- We know that:

$$f(A) = g(A) + h(A) = (\text{path cost to } A) + (\text{est. cost of } A \text{ to } A)$$

$$f(B) = g(B) + h(B) = (\text{path cost to } B) + (\text{est. cost of } B \text{ to } B)$$

- The heuristic must underestimate the true cost:

$$h(A) = h(B) = 0$$

- So now, we have to compare:

$$f(A) = g(A) = (\text{path cost to } A)$$


$$f(B) = g(B) = (\text{path cost to } B)$$

- We assumed that B is suboptimal! So

$$(\text{path cost to } A) < (\text{path cost to } B)$$

$$g(A) < g(B)$$


$$f(A) < f(B)$$



# Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B
- All ancestors of A expand before B
- A expands before B
- A\* search is optimal




# Properties of $A^*$

# Properties of $A^*$

Uniform-Cost





$A^*$



# UCS vs A\* Contours

- Uniform-cost expands equally in all “directions”
- A\* expands mainly toward the goal, but does hedge its bets to ensure optimality



[Demo: contours UCS / greedy / A\* empty (L3D1)]

[Demo: contours A\* pacman small maze (L3D5)]

# Video of Demo Contours (Empty) -- UCS

---



# Video of Demo Contours (Empty) -- Greedy

---





# Video of Demo Contours (Empty) – A\*

---



# Video of Demo Contours (Pacman Small Maze) – A\*

---



# Comparison



Greedy



Uniform Cost



A\*

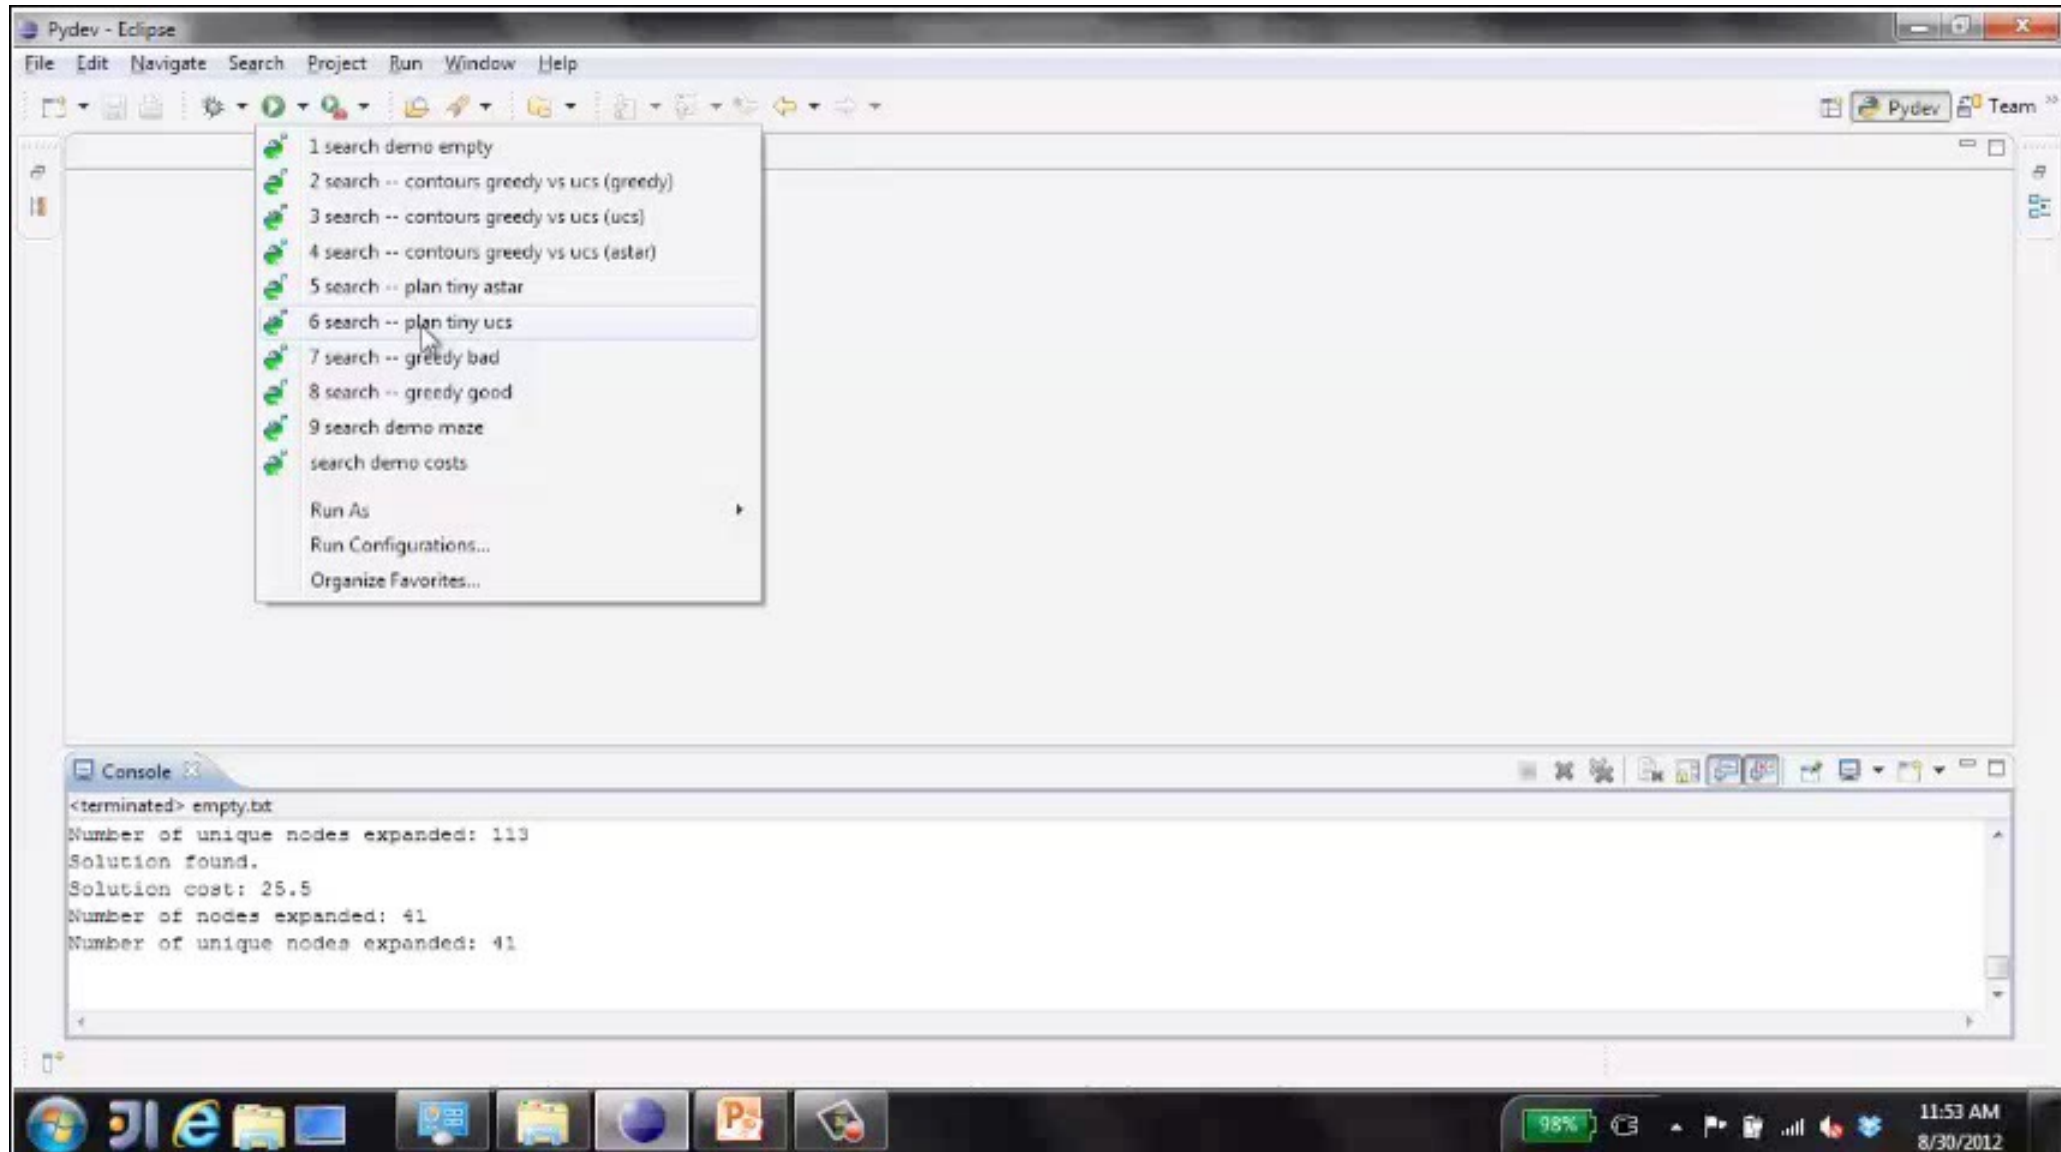
# A\* Applications

- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...

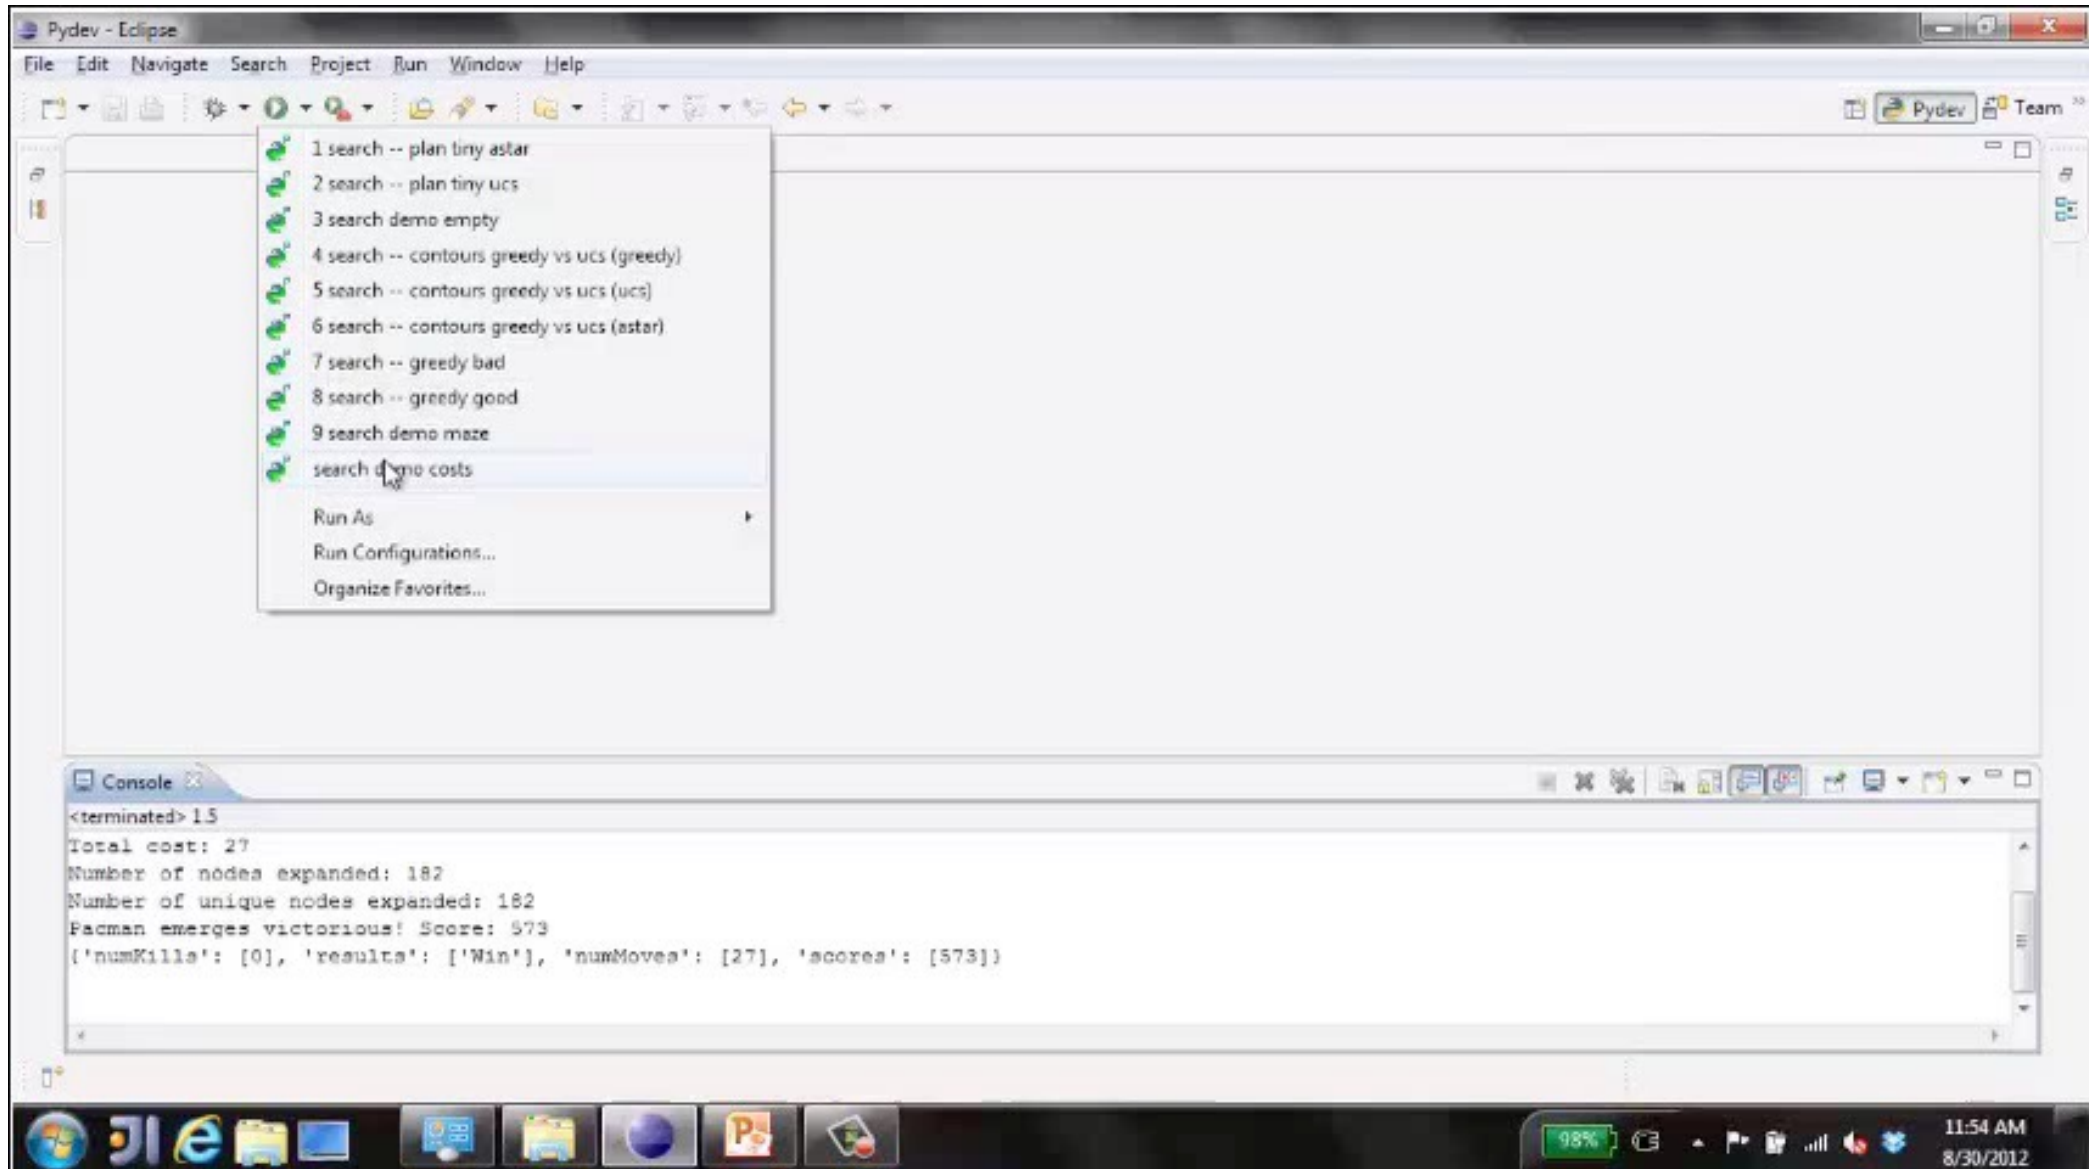


[Demo: UCS / A\* pacman tiny maze (L3D6,L3D7)]  
[Demo: guess algorithm Empty Shallow/Deep (L3D8)]

# Video of Demo Pacman (Tiny Maze) – UCS / A\*




# Video of Demo Empty Water Shallow/Deep – Guess Algorithm



# Creating Heuristics



---





# Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available




- Inadmissible heuristics are often useful too



# Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



Actions

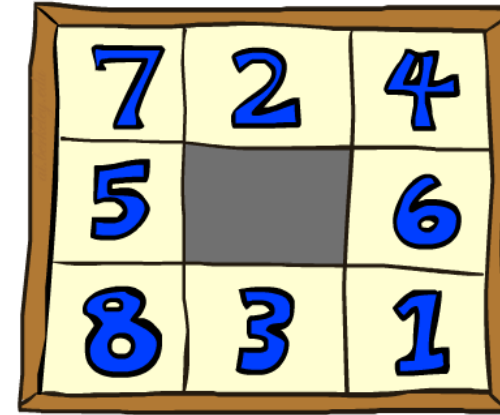
	1	2
3	4	5
6	7	8

Goal State

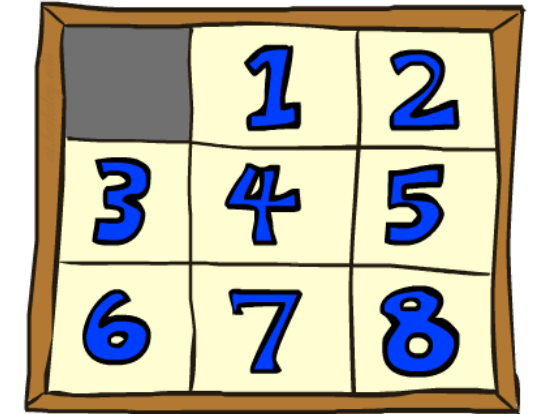
- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

# 8 Puzzle I


- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a *relaxed-problem* heuristic



Start State




Goal State



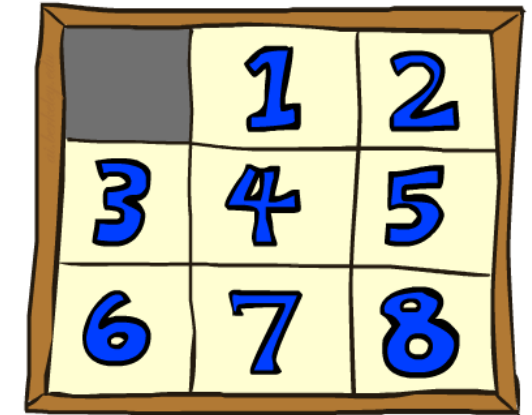
Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	$3.6 \times 10^6$
TILES	13	39	227

# 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance
- Why is it admissible?
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$



Start State




Goal State

Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTAN	12	25	73

# 8 Puzzle III

- How about using the *actual cost* as a heuristic?

- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?



- With  $A^*$ : a trade-off between quality of estimate and work per node

- As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

# Semi-Lattice of Heuristics

# Trivial Heuristics, Dominance


- Dominance:  $h_a \geq h_c$  if

$$\forall n : h_a(n) \geq h_c(n)$$


- Heuristics form a semi-lattice:
  - Max of admissible heuristics is admissible

$$h(n) = \max(h_a(n), h_b(n))$$

- Trivial heuristics
  - Bottom of lattice is the zero heuristic (what does this give us?)
  - Top of lattice is the exact heuristic





# Graph Search



# Tree Search: Extra Work!


- Failure to detect repeated states can cause exponentially more work.





# Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)




# Graph Search

---


- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states (“closed set”)
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list
- Can graph search wreck completeness? Why/why not?
- How about optimality?

# A\* Graph Search Gone Wrong?

State space graph



Search tree




Closed set


{ S B }

# A\* Graph Search Gone Wrong?

State space graph



Search tree




Closed set


{ S B C }

# A\* Graph Search Gone Wrong?

State space graph



Search tree




Closed set


{ S B C }

# A\* Graph Search Gone Wrong?

State space graph




Search tree



Closed set

{ S B C G }


# Consistency of Heuristics



- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal
$$h(A) \leq \text{actual cost from A to G}$$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc
$$h(A) - h(C) \leq \text{cost(A to C)}$$
- Consequences of consistency:
  - The f value along a path never decreases
$$h(A) \leq \text{cost(A to C)} + h(C)$$
  - A\* graph search is optimal

# Optimality of A\* Graph Search


---





# Optimality

- Tree search:
  - A\* is optimal if heuristic is admissible
  - UCS is a special case ( $h = 0$ )
- Graph search:
  - A\* optimal if heuristic is consistent
  - UCS optimal ( $h = 0$  is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems




# A\*: Summary

---




# A\*: Summary

- A\* uses both backward costs and (estimates of) forward costs
- A\* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems




# Search and Models

- Search operates over models of the world
  - The agent doesn't actually try all the plans out in the real world!
  - Planning is all “in simulation”
  - Your search is only as good as your models...







# Search Gone Wrong?



# Search Gone Wrong?



© 2005 MapQuest.com, Inc.



# Appendix: Search Pseudo-Code



# Tree Search Pseudo-Code

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
    end
  end
```



# Graph Search Pseudo-Code

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
  end
```