




Recap: Search Heuristics

- A heuristic is:
 - A function that *estimates* how close a state is to a goal
 - Designed for a particular search problem
 - Examples: Manhattan distance, Euclidean distance for pathing



Recap: Cost- vs. Heuristic-Guided Search



Uniform-Cost Search
(only costs, g)




Greedy Best-First Search
(only heuristic, h)




A* Search
(both, $f=g+h$)

Recap: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe




Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

Recap: 8-Puzzle

7	2	4
5		6
8	3	1


Start State



Actions


	1	2
3	4	5
6	7	8

Goal State




Designing a Heuristic: Knight's moves

- Minimum number of knight's moves to get from S to G?
 - $h_1 = (\text{Manhattan distance})/3$
 - $h_1' = h_1$ rounded up to correct parity (even if S, G same color, odd otherwise)
 - $h_2 = (\text{Euclidean distance})/\sqrt{5}$
 - $h_2' = h_2$ rounded up to correct parity
 - $h_3 = (\text{maximum horizontal or vertical distance})/2$
 - $h_3' = h_3$ rounded up to correct parity
- $h(n) = \max(h_1'(n), h_2'(n), h_3'(n))$ is admissible!





Recap: Optimality of A* Tree Search




Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.




Graph Search



Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)




Graph Search


- Idea: never **expand** a state twice
- How to implement:
 - Tree search + set of expanded states (“closed set”)
 - Expand the search tree node-by-node, but...
 - Before expanding a node, check to make sure its state has never been expanded before
 - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list
- Can graph search wreck completeness? Why/why not?
- How about optimality?

A* Graph Search Gone Wrong?

State space graph



Search tree




Closed set


{ S B }

A* Graph Search Gone Wrong?

State space graph



Search tree




Closed set


{ S B }

A* Graph Search Gone Wrong?

State space graph




Search tree



Closed set

{ S B C A }


Consistency of Heuristics



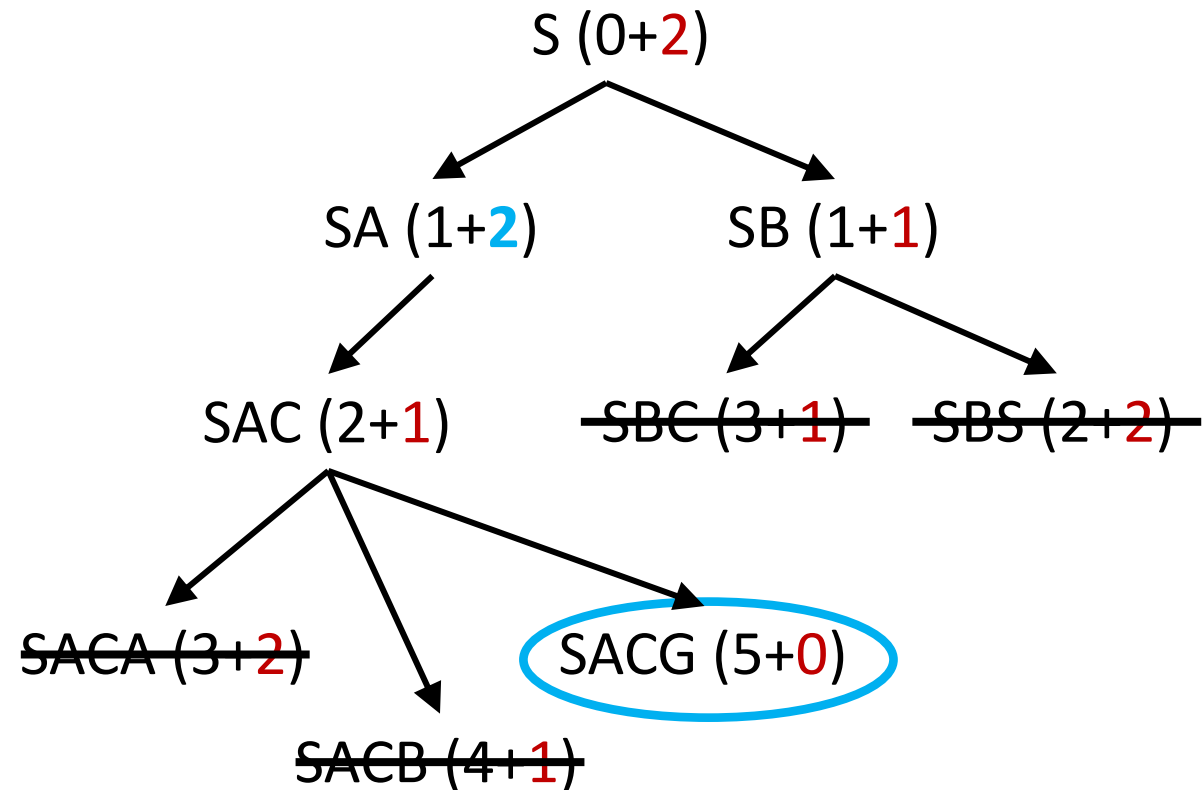
- Main idea: estimated heuristic costs \leq actual costs
 - Admissibility: heuristic cost \leq actual cost to goal
$$h(A) \leq \text{actual cost } h^* \text{ from A to G}$$
 - Consistency: heuristic “arc” cost \leq actual cost for each arc
$$h(A) - h(C) \leq \text{cost(A to C)}$$
 - a.k.a. “triangle inequality”: $h(A) \leq \text{cost(A to C)} + h(C)$
 - Note: true cost h^* necessarily satisfies triangle inequality
- Consequences of consistency:
 - The f value along a path never decreases
$$h(A) \leq \text{cost(A to C)} + h(C)$$
 - A* graph search is optimal

A* Graph Search with Consistent Heuristic

State space graph



Search tree

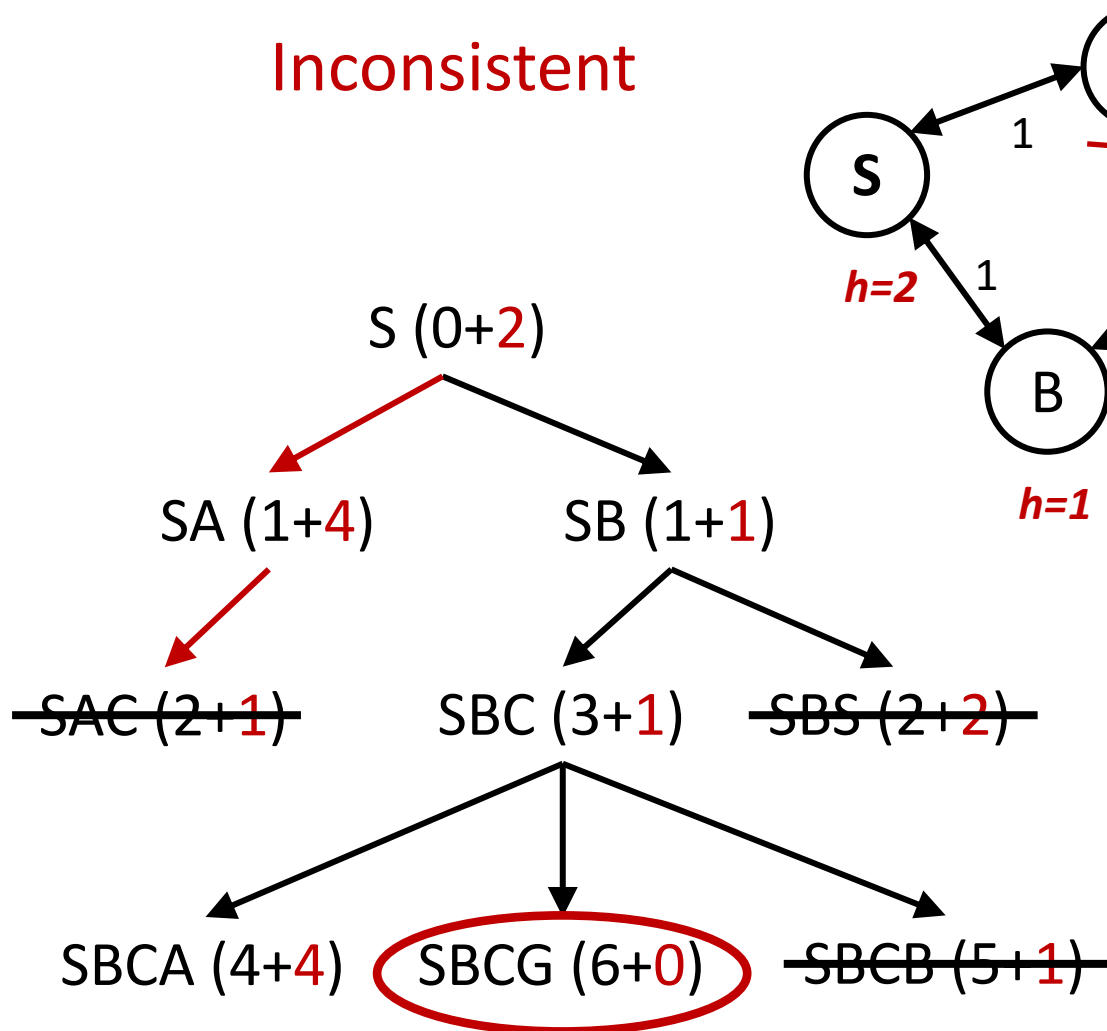


Closed set


{ S B A C }

Consistency => non-decreasing f-score

Inconsistent




Consistent




Optimality of A* Graph Search

- Sketch: consider what A* does with a consistent heuristic:
 - Fact 1: In tree search, A* expands nodes in increasing total f value (f-contours)
 - Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally
 - Result: A* graph search is optimal




Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is consistent
 - UCS optimal ($h = 0$ is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems




But...


- A* keeps the entire explored region in memory
- => will run out of space before you get bored waiting for the answer 
- There are variants that use less memory (Section 3.5.5):
 - IDA* works like iterative deepening, except it uses an f -limit instead of a depth limit
 - On each iteration, remember the smallest f -value that exceeds the current limit, use as new limit
 - Very inefficient when f is real-valued and each node has a unique value
 - RBFS is a recursive depth-first search that uses an f -limit = the f -value of the best alternative path available from any ancestor of the current node
 - When the limit is exceeded, the recursion unwinds but remembers the best reachable f -value on that branch
 - SMA* uses *all available memory* for the queue, minimizing thrashing
 - When full, drop worst node on the queue but remember its value in the parent

Search and Models


- Search operates over models of the world
 - The agent doesn't actually try all the plans out in the real world!
 - Planning is all “in simulation”
 - Your search is only as good as your models...




Search Gone Wrong?



Search Gone Wrong?



© 2005 MapQuest.com, Inc.



Tree Search Pseudo-Code

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
    end
  end
```

Graph Search Pseudo-Code

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
  end
```

Local Search





[These slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley. All CS188 materials are available at <http://ai.berkeley.edu>.]

[Updated slides from: Stuart Russell and Dawn Song]

Local search algorithms

- In many optimization problems, **path** is irrelevant; the goal state **is** the solution
- Then state space = set of “complete” configurations;
find **configuration satisfying constraints**, e.g., n-queens problem; or, find **optimal configuration**, e.g., travelling salesperson problem



- In such cases, can use **iterative improvement** algorithms: keep a single “current” state, try to improve it
- Constant space, suitable for online as well as offline search
- More or less unavoidable if the “state” is yourself (i.e., learning)


Hill Climbing

- Simple, general idea:
 - Start wherever
 - Repeat: move to the best neighboring state
 - If no neighbors better than current, quit





Heuristic for n -queens problem



- Goal: n queens on board with no **conflicts**, i.e., no queen attacking another
- States: n queens on board, one per column
- Actions: move a queen in its column
- Heuristic value function: number of conflicts



$h = 5$



$h = 2$




$h = 0$

Hill-climbing algorithm

```
function HILL-CLIMBING(problem) returns a state
  current ← make-node(problem.initial-state)
  loop do
    neighbor ← a highest-valued successor of current
    if neighbor.value ≤ current.value then
      return current.state
    current ← neighbor
```

“Like climbing Everest in thick fog with amnesia”

Global and local maxima



Random restarts


- find global optimum
- duh

Random sideways moves

- Escape from shoulders
- Loop forever on flat local maxima

Hill-climbing on the 8-queens problem

- **No sideways moves:**
 - Succeeds w/ prob. 0.14
 - Average number of moves per trial:
 - 4 when succeeding, 3 when getting stuck
 - Expected total number of moves needed:
 - $3(1-p)/p + 4 \approx 22$ moves
- **Allowing 100 sideways moves:**
 - Succeeds w/ prob. 0.94
 - Average number of moves per trial:
 - 21 when succeeding, 65 when getting stuck
 - Expected total number of moves needed:
 - $65(1-p)/p + 21 \approx 25$ moves



Moral: algorithms with knobs to twiddle are irritating

Simulated annealing

- Resembles the annealing process used to cool metals slowly to reach an ordered (low-energy) state
- Basic idea:
 - Allow “bad” moves occasionally, depending on “temperature”
 - High temperature => more bad moves allowed, shake the system out of its local minimum
 - Gradually reduce temperature according to some schedule
 - Sounds pretty flaky, doesn't it?

Simulated annealing algorithm

function SIMULATED-ANNEALING(problem,schedule) **returns** a state

current \leftarrow problem.initial-state

for t = 1 **to** ∞ **do**

 T \leftarrow schedule(t)


if T = 0 **then return** current

 next \leftarrow a randomly selected successor of current

$\Delta E \leftarrow$ next.value – current.value



if $\Delta E > 0$ **then** current \leftarrow next

else current \leftarrow next only with probability $e^{\Delta E/T}$




Simulated Annealing

- Theoretical guarantee:
 - Stationary distribution (Boltzmann): $P(x) \propto e^{E(x)/T}$
 - If T decreased slowly enough, will converge to optimal state!
- Proof sketch
 - Consider two adjacent states x, y with $E(y) > E(x)$ [high is good]
 - Assume $x \rightarrow y$ and $y \rightarrow x$ and outdegrees $D(x) = D(y) = D$
 - Let $P(x), P(y)$ be the equilibrium occupancy probabilities at T
 - Let $P(x \rightarrow y)$ be the probability that state x transitions to state y




Occupation probability as a function of T



Simulated Annealing

- Is this convergence an interesting guarantee?
- Sounds like magic, but reality is reality:
 - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
 - “Slowly enough” may mean exponentially slowly
 - Random restart hillclimbing also converges to optimal state...
- Simulated annealing and its relatives are a key workhorse in VLSI layout and other optimal configuration problems




Local beam search

- Basic idea:
 - K copies of a local search algorithm, initialized randomly
 - For each iteration
 - Generate ALL successors from K current states
 - Choose best K of these to be the new current states




Or, K chosen randomly with
a bias towards good ones

Beam search example ($K=4$)





Local beam search

- Why is this different from K local searches in parallel?
 - The searches **communicate**! “Come over here, the grass is greener!”
- What other well-known algorithm does this remind you of?
 - Evolution!




Genetic algorithms



- Genetic algorithms use a natural selection metaphor
 - Resample *K* individuals at each step (selection) weighted by fitness function
 - Combine by pairwise crossover operators, plus mutation to give variety

Example: N-Queens




- Does crossover make sense here?
- What would mutation be?
- What would a good fitness function be?

Local search in continuous spaces



Example: Placing airports in Romania

Place 3 airports to minimize the sum of squared distances from each city to its nearest airport



Airport locations

$$\mathbf{x} = (x_1, y_1), (x_2, y_2), (x_3, y_3)$$

City locations (x_c, y_c)

C_a = cities closest to airport a

Objective: minimize

$$f(\mathbf{x}) = \sum_a \sum_{c \in C_a} (x_a - x_c)^2 + (y_a - y_c)^2$$

Handling a continuous state/action space

1. Discretize it!

- Define a grid with increment δ , use any of the discrete algorithms

2. Choose random perturbations to the state

- First-choice hill-climbing: keep trying until something improves the state
- Simulated annealing

3. Compute gradient of $f(\mathbf{x})$ analytically

Finding extrema in continuous space

- Gradient vector $\nabla f(\mathbf{x}) = (\partial f / \partial x_1, \partial f / \partial y_1, \partial f / \partial x_2, \dots)^\top$
- For the airports, $f(\mathbf{x}) = \sum_a \sum_{c \in C_a} (x_a - x_c)^2 + (y_a - y_c)^2$
- $\partial f / \partial x_1 = \sum_{c \in C_1} 2(x_1 - x_c)$
- At an extremum, $\nabla f(\mathbf{x}) = 0$
- Can sometimes solve in closed form: $x_1 = (\sum_{c \in C_1} x_c) / |C_1|$
 - Is this a local or global minimum of f ?
- If we can't solve $\nabla f(\mathbf{x}) = 0$ in closed form...
 - Gradient descent: $\mathbf{x} \leftarrow \mathbf{x} - \alpha \nabla f(\mathbf{x})$
- Huge range of algorithms for finding extrema using gradients

Summary

- Many configuration and optimization problems can be formulated as local search
- General families of algorithms:
 - Hill-climbing, continuous optimization
 - Simulated annealing (and other stochastic methods)
 - Local beam search: multiple interaction searches
 - Genetic algorithms: break and recombine states

Many machine learning algorithms are local searches